

Vitaliy Kuzmin

Senior Agentic AI Engineer | Tbilisi, Georgia | Remote / Full-time / Contract

viton930@gmail.com | linkedin.com/in/vitaliy-kuzmin-ai-engineer | vitalykuzmin.github.io/portfolio

Career Profile

Senior Agentic AI Engineer and Full Stack Engineer with 8+ years of experience building complex software systems and 3+ years focused on LLM orchestration, RAG, AI agents, local/on-premise models, and enterprise AI platforms.

Experienced in building practical AI systems around real business workflows: multi-agent orchestration, private AI infrastructure, domain RAG, SQL agents, LLM evaluation, local LLM deployment, observability, enterprise API integrations, and technical leadership.

Career Highlights

- Reduced manual navigation across enterprise data and documentation by up to 50% through an enterprise AI assistant platform.
- Reduced repetitive editorial preparation and formatting work by up to 60% through an AI-assisted publishing workflow.
- Led engineering delivery for a 10+ person cross-functional team across backend, C++, QA, design, analytics, and product work.

Skills

- **AI & LLM:** LangGraph, LangChain, LlamaIndex, RAG, SQL agents, multi-agent systems, MCP, local LLMs, Ollama, LiteLLM, Langfuse, LLM-as-a-judge, prompt engineering.
- **Backend & Infrastructure:** Python, FastAPI, REST, gRPC, Docker, Google Cloud Platform (GCP), CI/CD, DevOps, PostgreSQL, Redis, vector databases, LanceDB, FastEmbed, OpenTelemetry.
- **Leadership:** architecture, technical leadership, team leadership, hiring, mentoring, product planning, Agile, Scrum, Kanban.

Professional Experience

InterProCom | Lead AI Architect | Feb 2025 - Present

InterProCom is a leading developer of Enterprise Asset Management (EAM) platforms and an IBM Premier Business Partner, building proprietary software for complex industrial workflows.

- Architected Axioma AI, an enterprise AI assistant platform for asset management workflows in the energy domain.
- Built LangGraph-based multi-agent workflows for SQL, norm-control, and domain RAG over enterprise data and documentation.
- Designed LLM-as-a-judge and automated evaluation workflows with Langfuse for agent quality control.
- Integrated AI workflows with internal APIs, structured data, documentation, model routing, observability, and private/on-premise LLM infrastructure.
- Organized Docker-based deployment, CI/CD, DevOps processes, and context-driven development with architecture docs, global prompts, agent instructions, reusable skills, and project knowledge alongside the codebase.
- Reduced manual navigation across enterprise data and documentation by up to 50% in target workflows and expanded AI-assisted analytics with broader enterprise data sources.

BL Group R&D Center | Lead AI Engineer | Jun 2022 - Feb 2025

BL Group R&D Center was an internal R&D unit focused on software prototypes, automation products, and AI-assisted workflows.

- Led Smart Publish, an AI-assisted publishing platform for creating, adapting, reviewing, scheduling, and publishing content across websites and social platforms.
- Designed a predictable LLM workflow for draft processing, rewriting, editing, formatting, metadata preparation, channel-specific adaptation, and human review.
- Integrated draft intake, Telegram workflows, CMS publishing through APIs, and preparation for social media distribution.
- Contributed to MVP/product experiments before Smart Publish, including cloud and automation workflows.
- Reduced repetitive editorial preparation and formatting work by up to 60% and improved content consistency through platform-specific adaptation.

BL Group | Team Lead / Lead Software Engineer | Jul 2018 - Jun 2022

BL Group built engineering software products for lighting design, lighting calculations, and urban lighting workflows.

- Grew from Full Stack Developer to Team Lead, managing cross-functional engineering work across C++, Python, frontend, QA, design, analytics, and product planning.
- Led development processes for Light in Night, a professional lighting calculation and design application.
- Managed and coordinated a 10+ person engineering team across product development, QA, releases, and technical planning.
- Led Urban Lighting, an ML reconstruction platform combining field inventory, mobile laboratory measurements, geospatial mapping, and LightGBM-based illuminance prediction.
- Accelerated field data collection and processing by up to 60%, reduced repeat city survey visits by up to 70%, and removed dependency on an external map-markup contractor.

Selected Projects

- **Axioma AI**: enterprise AI assistant platform using LangGraph, RAG, SQL agents, local LLM infrastructure, Langfuse evaluation, observability, and private/on-premise deployment.
- **hAi**: local-first desktop AI assistant and personal R&D environment for AI agents, memory, voice interaction, local models, MCP, LLM evaluation, and desktop context.
- **Smart Publish**: AI-assisted publishing workflow for draft intake, content generation, channel-specific adaptation, review, scheduling, CMS publishing, GCP, and social distribution.
- **Urban Lighting**: ML reconstruction platform for field inventory, mobile laboratory measurements, geospatial datasets, LightGBM-based prediction, and planning workflows.

Education

Moscow Power Engineering Institute (MPEI) - Master's Degree, Electronics and Nanoelectronics, 2018